

Statistical Computing

Michael Mayer

March 2023



Statistical Computing: What will we do?

Chapters

1. R in Action
2. Statistical Inference
3. Linear Models
4. Model Selection and Validation
5. Trees
6. Neural Nets

Remarks

- ▶ Chapters 3 to 6:
Statistical ML in Action
- ▶ Two weeks per chapter
- ▶ Exercises at end of chapter notes

Trees

Outline

- ▶ Decision Trees
- ▶ Random Forests
- ▶ Gradient Boosted Trees

Decision Trees

- ▶ Simple
- ▶ Easy to interpret
- ▶ Decision trees are like wolves:
Weak alone, strong together
- ▶ Around since 1984
(Breiman, Friedman)



<https://images.pexels.com/photos/3732527/pexels-photo-3732527.jpeg>

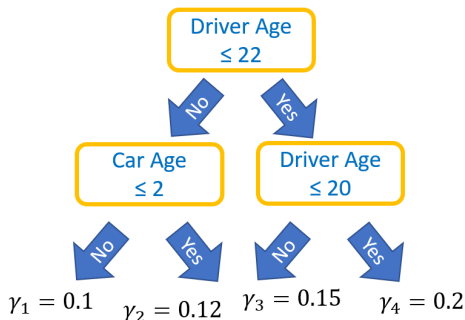
What is a Decision Tree?

Greedy recursive partitioning

1. Split: find best “yes/no” question on best feature to make total loss smaller
2. Apply Step 1 recursively

Predictions

- ▶ Follow splits and use leaf value γ_j
- ▶ Usually, γ_j is average response in leaf j
- ▶ Terminal regions R_1, \dots, R_J
- ▶ \mathbf{x} falls in leaf $j \Leftrightarrow \mathbf{x} \in R_j$
- ▶ $\hat{f}(\mathbf{x}) = \sum_{j=1}^J \gamma_j \mathbf{1}\{\mathbf{x} \in R_j\}$



The tree does a headstand

Example

Properties of Decision Trees

- ▶ Outliers
- ▶ Missing values
- ▶ Categorical covariates
- ▶ Greedy
- ▶ Interactions
- ▶ Extrapolation
- ▶ Instability

Most properties are inherited to groups/ensembles of decision trees

From nearest neighbors to decision trees: Short video by Jerome Friedman:
<https://www.youtube.com/watch?v=8hupHmBVvb0>

Random Forests

- ▶ Combine many decision trees
- ▶ Perform very well
- ▶ Black Box
- ▶ Around since 2001 (Breiman)
- ▶ Why is combination of trees better than a single one?



<https://images.pexels.com/photos/1459534/pexels-photo-1459534.jpeg>

Ensembling and Bagging

Ensembling

- ▶ Combine multiple models (**base learners**) to single one
- ▶ Example: k -nearest-neighbor with different k
- ▶ Combined predictions have lower variance \rightarrow better test performance (diversified stock portfolio, Bias-Variance Trade-Off)

Algorithm: Bagging (**B**ootstrap **a**ggregating)

1. Select B bootstrapped training data sets from the original training data
2. Fit model $\hat{f}^{*j}(\mathbf{x})$ on each of them
3. Return the bagged model $\hat{f}(\mathbf{x}) = \frac{1}{B} \sum_{j=1}^B f^{*j}(\mathbf{x})$

Example

Bias-Variance Trade-Off

Decomposition of generalization error:

$$\underbrace{\mathbb{E}(y_o - \hat{f}(x_o))^2}_{\text{Expected test MSE of } x_o} = \text{Var}(\hat{f}(x_o)) + [\text{Bias}(\hat{f}(x_o))]^2 + \underbrace{\text{Var}(\varepsilon)}_{\text{Irreducible error}}$$

- ▶ One specific observation (y_o, x_o)
- ▶ Expectations and variances over large number of training sets
- ▶ Bias: Error introduced by approximating true model by f
- ▶ Low bias \leftrightarrow high variance \rightarrow Bias-Variance Trade-Off
- ▶ Bagged decision trees: low bias (why?) and low variance

Remarks on Bagging

- ▶ Works best with unstable base learners → deep decision trees
- ▶ Out-of-bag (OOB) validation
- ▶ Parallel computing
- ▶ Performance versus complexity

From Bagging to Random Forests

A random forest is a bagged decision tree with an extra twist

Twist

- ▶ Additional source of randomness
- ▶ Each split considers only random feature subset (often $p/3$ or \sqrt{p})
- ▶ Additional decorrelation \rightarrow stronger diversification

Algorithm: Random forest (regression)

1. Select B bootstrapped training data sets from the original training data
2. Fit (usually deep) decision tree $\hat{f}^{*j}(\mathbf{x})$ on each of them. For each split, consider only random feature subset
3. Return the random forest $\hat{f}(\mathbf{x}) = \frac{1}{B} \sum_{j=1}^B \hat{f}^{*j}(\mathbf{x})$

Comments on Random Forests

- ▶ Number of trees
- ▶ Deep trees
- ▶ Don't trust performance on training set → OOB performance
- ▶ Parameter tuning

Regarding **parameter tuning**: Short video by Adele Cutler on working with Leo Breiman: https://www.youtube.com/watch?v=t8ooi_tJHSE

Example

Interpreting a Black Box

Study

1. Performance
2. Variable importance
3. Effects

XAI

- ▶ e**X**plainable **A**rtificial **I**ntelligence
- ▶ Collection of methods to interpret models
- ▶ Examples: Split-gain importance, ICE, PDP

Example

- ▶ Split gain importance of random forest
- ▶ Variable importance and linear regression?

Individual Conditional Expectation (ICE)

Basic thinking

- ▶ In **additive** linear model f , the effect of $X^{(j)}$ is fully described by its coefficient(s)
- ▶ It describes how f reacts on changes in $X^{(j)}$ (Ceteris Paribus)
- ▶ What if model involves complex interactions?

Idea (Goldstein et al., 2015)

- ▶ Study (Ceteris Paribus) effect of $X^{(j)}$ for **one** observation
- ▶ *ICE function* for feature $X^{(j)}$ of model f and observation $\mathbf{x} \in \mathbb{R}^p$

$$\text{ICE}_j : v \in \mathbb{R} \mapsto f(v, \mathbf{x}_{\setminus j})$$

- ▶ $\mathbf{x}_{\setminus j}$ denotes all but the j -th component of \mathbf{x} , which is replaced by v
- ▶ *ICE curve* represents graph $(v, \text{ICE}_j(v))$ for grid of values $v \in \mathbb{R}$

ICE Plot: Visualize ICE Curves of many Observations

Example

Notes

- ▶ Curves with different shapes indicate interaction effects
- ▶ Parallel curves \Leftrightarrow additivity in $X^{(j)}$
- ▶ Centered ICE plots
- ▶ Usually on link scale (why?)

Pros and Cons

- + Simple to compute
- + Easy to interpret (Ceteris Paribus)
- + Gives impression about interactions
- Ceteris Paribus can be unnatural
- Model applied to rare/impossible \mathbf{x}

Partial Dependence Plot PDP (Friedman 2001)

- ▶ Average of many ICE curves
- ▶ Ceteris Paribus effect of $X^{(j)}$ averaged over all interaction effects
- ▶ (Empirical) partial dependence function of j -th feature

$$\text{PD}_j(v) = \frac{1}{n} \sum_{i=1}^n \hat{f}(v, \mathbf{x}_{i,\setminus j})$$

- ▶ $\mathbf{x}_{i,\setminus j}$ feature vector of i -th observation without j -th component
- ▶ PDP equals graph $(v, \text{PD}_j(v))$ for grid of values $v \in \mathbb{R}$
- ▶ Sum runs over reference data (=?)
- ▶ Pros/cons similar to ICE, but no info on interaction

Example

Gradient Boosted Trees

- ▶ Combine many decision trees
- ▶ Perform very well
- ▶ Black Box
- ▶ Around since 2001 (Friedman)

} Like random forests



<https://www.gormananalysis.com/blog/gradient-boosting-explained/>

Boosting

Basic idea of boosting (e.g. Schapire, 1990)

1. Fit simple model \hat{f} to data
2. For $k = 1, \dots, K$ do:
 - a. Find simple model \hat{f}^k that corrects the mistakes of \hat{f}
 - b. Update: $\hat{f} \leftarrow \hat{f} + \hat{f}^k$

How to find updates \hat{f}^k ?

Use decision trees \rightarrow **boosted trees**

- ▶ Use reweighting heuristic for binary classification
 \rightarrow AdaBoost (Freund and Schapire, 1995)
- ▶ Reduce total loss $Q(\hat{f} + \hat{f}^k) = \sum_{i=1}^n L(y_i, \hat{f}(\mathbf{x}_i) + \hat{f}^k(\mathbf{x}_i))$
 \rightarrow **Gradient boosting** (Friedman, 2001)

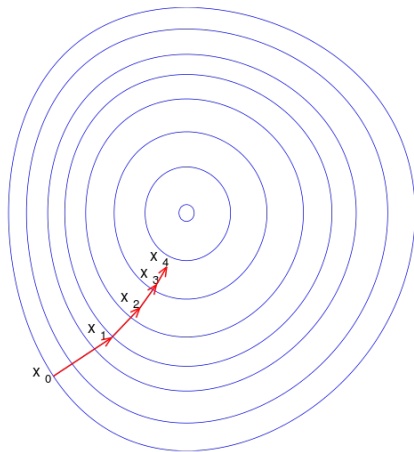
Gradient Descent

Minimize function $h : \mathbb{R}^n \rightarrow \mathbb{R}$

1. Start at some value $\hat{x} \in \mathbb{R}^n$
2. Repeat: $\hat{x} \leftarrow \hat{x} - \lambda g$
 - ▶ $\lambda > 0$: Step size or learning rate
 - ▶ Gradient $g \in \mathbb{R}^n$ of h at \hat{x} :

$$g = \left[\frac{\partial h(x)}{\partial x} \right]_{x=\hat{x}}$$

- ▶ g points in direction of steepest ascent



https://en.wikipedia.org/wiki/Gradient_descent

Gradient Boosting

Gradient descent of $Q(f)$

- ▶ $Q(f) = \sum_{i=1}^n L(y_i, f(\mathbf{x}_i))$
- ▶ $f = (f(\mathbf{x}_1), \dots, f(\mathbf{x}_n)) \in \mathbb{R}^n$

1. Start at some value \hat{f}
2. Repeat: $\hat{f} \leftarrow \hat{f} - \lambda g$ with

$$g = \left[\frac{\partial Q(f)}{\partial f} \right]_{f=\hat{f}}$$

having components

$$g_i = \left[\frac{\partial L(y_i, f(\mathbf{x}_i))}{\partial f(\mathbf{x}_i)} \right]_{f(\mathbf{x}_i)=\hat{f}(\mathbf{x}_i)}$$

For squared error?

- ▶ $L(y, z) = (y - z)^2/2$
- ▶ Plugging in: $g_i = -\underbrace{(y_i - \hat{f}(\mathbf{x}_i))}_{\text{Residual } r_i}$
- ▶ Repeat: $\hat{f}(\mathbf{x}_i) \leftarrow \hat{f}(\mathbf{x}_i) + \underbrace{\lambda r_i}_{\hat{f}^k?}$

Boosting with $\hat{f}^k = -\lambda g_i$? Now way...

1. y_i unknown in application
 2. Should work for all \mathbf{x}
- replace $-g_i$ by predictions of tree

Gradient Boosted Trees for Squared Error Loss

Algorithm

1. Initialize $\hat{f}(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n y_i$
2. For $k = 1, \dots, K$ do:
 - a. For $i = 1, \dots, n$, calculate residuals $r_i = y_i - \hat{f}(\mathbf{x}_i)$
 - b. Model the r_i as a function of the \mathbf{x}_i by fitting a regression tree \hat{f}^k
 - c. Update: $\hat{f}(\mathbf{x}) \leftarrow \hat{f}(\mathbf{x}) + \lambda \hat{f}^k(\mathbf{x})$
3. Output $\hat{f}(\mathbf{x})$

General loss functions?

- ▶ Replace residuals by negative gradients of loss function
- ▶ Leaf values might be suboptimal \rightarrow replace by optimal values

Gradient Boosted Trees for General Losses

1. Initialize $\hat{f}(\mathbf{x}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$
2. For $k = 1, \dots, K$ do:
 - a. For $i = 1, \dots, n$, calculate negative gradients (pseudo-residuals)

$$r_i = - \left[\frac{\partial L(y_i, f(\mathbf{x}_i))}{\partial f(\mathbf{x}_i)} \right]_{f(\mathbf{x}_i) = \hat{f}(\mathbf{x}_i)}$$

- b. Model r_i as function of \mathbf{x}_i by regression tree \hat{f}^k with terminal regions R_1, \dots, R_J
- c. For each $j = 1, \dots, J$, use line-search to find the optimal leaf value

$$\gamma_j = \operatorname{argmin}_{\gamma} \sum_{\mathbf{x}_i \in R_j} L(y_i, \hat{f}(\mathbf{x}_i) + \gamma)$$

- d. Update: $\hat{f}(\mathbf{x}) \leftarrow \hat{f}(\mathbf{x}) + \underbrace{\lambda \sum_{j=1}^J \gamma_j \mathbf{1}\{\mathbf{x} \in R_j\}}_{\text{modified tree}}$

3. Output $\hat{f}(\mathbf{x})$

Remarks

- ▶ Predictions are sum of short decision trees (with modified leaf values)
- ▶ Random forest: average of deep trees
- ▶ How to select learning rate λ , number of trees K , ...?
- ▶ (AdaBoost is gradient tree boosting with exponential loss)

Modern Implementations

Timeline

1. XGBoost (2014)
2. LightGBM (2016)
3. CatBoost (2017)

Example

Differences to Friedman's original

- ▶ Use of second order gradients
→ no line-search necessary
- ▶ Histogram binning
→ speeds up tree growth
- ▶ Penalized objective function

Parameter Tuning is Essential

1. Number of boosting rounds/trees K
→ find by early stopping (validation/CV)
2. Learning rate λ
→ to get reasonable number of rounds
3. Regularization
 - ▶ Tree depth, number of leaves, loss penalties, etc.
 - ▶ → Grid/Randomized search and iterate process

Example

- ▶ XGBoost
- ▶ LightGBM

Comments

- ▶ Why not one big grid search on all parameters?
- ▶ Objective/metrics